

car-price-prediction-randomforestregression-vs-xgbregression

February 27, 2024

1 CAR PRICE PREDICTION & EDA WITH XGBoost Regression

2 Import library

```
[1]: #Linear algebra & data processing
import numpy as np
import pandas as pd

#Data visualization libraries
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

#Import transofmers
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder

#Import Regression method
from sklearn.svm import SVR
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.linear_model import Ridge, ElasticNet, Lasso, LogisticRegression,   
↳ LinearRegression

#Import model selection
from sklearn.model_selection import train_test_split, cross_val_score,   
↳ GridSearchCV, KFold

#Import Accuracy Metrics
from sklearn.metrics import r2_score, max_error, mean_squared_error,   
↳ mean_absolute_error
from time import time
```

```
import warnings
warnings.filterwarnings("ignore")
```

3 Some functions

```
[2]: def pourcentage(data):
      n = data.shape[0]
      ret = pd.DataFrame(data.isnull().sum(), columns=['missing_number'])
      ret['pourcentage_missing_number'] = (ret['missing_number']/n)*100
      ret['types'] = data.dtypes
      ret['duplicate'] = data.duplicated(keep=False).sum()
      ret['NAN'] = data.isna().sum()
      return ret
```

4 Import data

```
[3]: df_car = pd.read_csv("/kaggle/input/carr-details/Car details v3.csv")
      df_car.head()
```

```
[3]:
```

		name	year	selling_price	km_driven	fuel	\
0		Maruti Swift Dzire VDI	2014	450000	145500	Diesel	
1		Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	
2		Honda City 2017-2020 EXi	2006	158000	140000	Petrol	
3		Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	
4		Maruti Swift VXi BSIII	2007	130000	120000	Petrol	

	seller_type	transmission	owner	mileage	engine	max_power	\
0	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	
1	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	
2	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	
3	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	
4	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	

	torque	seats
0	190Nm@ 2000rpm	5.0
1	250Nm@ 1500-2500rpm	5.0
2	12.7@ 2,700(kgm@ rpm)	5.0
3	22.4 kgm at 1750-2750rpm	5.0
4	11.5@ 4,500(kgm@ rpm)	5.0

5 Pre-processing Data

```
[4]: pourcentage(df_car)
```

```
[4]:
```

	missing_number	pourcentage_missing_number	types	duplicate	\
name	0	0.000000	object	1827	
year	0	0.000000	int64	1827	
selling_price	0	0.000000	int64	1827	
km_driven	0	0.000000	int64	1827	
fuel	0	0.000000	object	1827	
seller_type	0	0.000000	object	1827	
transmission	0	0.000000	object	1827	
owner	0	0.000000	object	1827	
mileage	221	2.718996	object	1827	
engine	221	2.718996	object	1827	
max_power	215	2.645177	object	1827	
torque	222	2.731299	object	1827	
seats	221	2.718996	float64	1827	

```

NAN
name      0
year      0
selling_price  0
km_driven  0
fuel      0
seller_type  0
transmission  0
owner      0
mileage    221
engine     221
max_power  215
torque     222
seats      221

```

```
[5]: df_car = df_car.dropna(axis=0)

def convertToNumber(s:str):
    d=""
    for i in list(s):
        if i.isdigit():
            d += i
    return eval(d)

df_car["mileage"] = df_car["mileage"].apply(convertToNumber)
df_car["engine"] = df_car["engine"].apply(convertToNumber)
df_car["max_power"] = df_car["max_power"].apply(convertToNumber)
```

```
[6]: pourcentage(df_car)
```

```
[6]:
```

	missing_number	pourcentage_missing_number	types	duplicate	\
name	0	0.0	object	1801	

year	0	0.0	int64	1801
selling_price	0	0.0	int64	1801
km_driven	0	0.0	int64	1801
fuel	0	0.0	object	1801
seller_type	0	0.0	object	1801
transmission	0	0.0	object	1801
owner	0	0.0	object	1801
mileage	0	0.0	int64	1801
engine	0	0.0	int64	1801
max_power	0	0.0	int64	1801
torque	0	0.0	object	1801
seats	0	0.0	float64	1801

	NAN
name	0
year	0
selling_price	0
km_driven	0
fuel	0
seller_type	0
transmission	0
owner	0
mileage	0
engine	0
max_power	0
torque	0
seats	0

```
[7]: df_car.head()
```

```
[7]:
```

		name	year	selling_price	km_driven	fuel	\
0		Maruti Swift Dzire VDI	2014	450000	145500	Diesel	
1		Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	
2		Honda City 2017-2020 EXi	2006	158000	140000	Petrol	
3		Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	
4		Maruti Swift VXi BSIII	2007	130000	120000	Petrol	

	seller_type	transmission	owner	mileage	engine	max_power	\
0	Individual	Manual	First Owner	234	1248	74	
1	Individual	Manual	Second Owner	2114	1498	10352	
2	Individual	Manual	Third Owner	177	1497	78	
3	Individual	Manual	First Owner	230	1396	90	
4	Individual	Manual	First Owner	161	1298	882	

		torque	seats
0		190Nm@ 2000rpm	5.0
1		250Nm@ 1500-2500rpm	5.0

```

2      12.7@ 2,700(kgm@ rpm)      5.0
3      22.4 kgm at 1750-2750rpm    5.0
4      11.5@ 4,500(kgm@ rpm)      5.0

```

```
[8]: data = df_car.drop(['name', 'torque', 'seller_type', 'owner'], axis=1)
data.head()
```

```
[8]:
```

	year	selling_price	km_driven	fuel	transmission	mileage	engine \
0	2014	450000	145500	Diesel	Manual	234	1248
1	2014	370000	120000	Diesel	Manual	2114	1498
2	2006	158000	140000	Petrol	Manual	177	1497
3	2010	225000	127000	Diesel	Manual	230	1396
4	2007	130000	120000	Petrol	Manual	161	1298

	max_power	seats
0	74	5.0
1	10352	5.0
2	78	5.0
3	90	5.0
4	882	5.0

```
[9]: data.describe()
```

```
[9]:
```

	year	selling_price	km_driven	mileage	engine \
count	7906.000000	7.906000e+03	7.906000e+03	7906.000000	7906.000000
mean	2013.983936	6.498137e+05	6.918866e+04	947.702378	1458.708829
std	3.863695	8.135827e+05	5.679230e+04	925.336832	503.893057
min	1994.000000	2.999900e+04	1.000000e+00	0.000000	624.000000
25%	2012.000000	2.700000e+05	3.500000e+04	185.000000	1197.000000
50%	2015.000000	4.500000e+05	6.000000e+04	240.000000	1248.000000
75%	2017.000000	6.900000e+05	9.542500e+04	1944.000000	1582.000000
max	2020.000000	1.000000e+07	2.360457e+06	3344.000000	3604.000000

	max_power	seats
count	7906.000000	7906.000000
mean	2766.125348	5.416393
std	5162.123778	0.959208
min	35.000000	2.000000
25%	100.000000	5.000000
50%	739.000000	5.000000
75%	3748.000000	5.000000
max	108495.000000	14.000000

```
[10]: data.fuel.unique()
```

```
[10]: array(['Diesel', 'Petrol', 'LPG', 'CNG'], dtype=object)
```

```
[11]: data_new = pd.get_dummies(data=data, columns=['fuel'], drop_first=True,
↳ dtype=int)
```

```
[12]: data_new.head()
```

```
[12]:   year  selling_price  km_driven  transmission  mileage  engine  max_power \
0  2014         450000    145500         Manual      234    1248         74
1  2014         370000    120000         Manual    2114    1498        10352
2  2006         158000    140000         Manual     177    1497         78
3  2010         225000    127000         Manual     230    1396         90
4  2007         130000    120000         Manual     161    1298         882

   seats  fuel_Diesel  fuel_LPG  fuel_Petrol
0     5.0           1         0           0
1     5.0           1         0           0
2     5.0           0         0           1
3     5.0           1         0           0
4     5.0           0         0           1
```

```
[13]: data_new["transmission"] = data_new["transmission"].replace({'Automatic': 1,
↳ 'Manual': 0})
```

```
[14]: data_new.head()
```

```
[14]:   year  selling_price  km_driven  transmission  mileage  engine  max_power \
0  2014         450000    145500           0      234    1248         74
1  2014         370000    120000           0    2114    1498        10352
2  2006         158000    140000           0     177    1497         78
3  2010         225000    127000           0     230    1396         90
4  2007         130000    120000           0     161    1298         882

   seats  fuel_Diesel  fuel_LPG  fuel_Petrol
0     5.0           1         0           0
1     5.0           1         0           0
2     5.0           0         0           1
3     5.0           1         0           0
4     5.0           0         0           1
```

6 Scaling data

```
[15]: mmScaler = MinMaxScaler()
mmScaler_y = MinMaxScaler()

label_enc = LabelEncoder()

x = data_new[['year', 'km_driven', 'transmission', 'mileage', 'engine',
↳ 'max_power', 'seats', 'fuel_Diesel', 'fuel_LPG', 'fuel_Petrol']].values
```

```
y = data_new[['selling_price']].values
```

```
[16]: x[:, 0] = label_enc.fit_transform(x[:, 0])
      x = mmScaler.fit_transform(x)
      y = mmScaler_y.fit_transform(y)
```


7 Exploration Data Analysis And Correlation

```
[17]: #Pair plot labels
      sns.pairplot(data_new)
```

```
[17]: <seaborn.axisgrid.PairGrid at 0x7f4458c5ba30>
```



```
[18]: sns.distplot(data_new['selling_price'], label="reals values", color='red')
plt.legend()
plt.grid()
plt.show()
```



```
[19]: #correlation
corr = pd.DataFrame(data_new.corrwith(data_new['selling_price']))
corr
```

```
[19]:
```

	0
year	0.412302
selling_price	1.000000
km_driven	-0.222158
transmission	0.590269
mileage	0.098988
engine	0.455682
max_power	0.137042
seats	0.041617
fuel_Diesel	0.204831
fuel_LPG	-0.035978


```
fuel_Petrol    -0.195074
```

8 Split Data

```
[20]: X_train, X_test, Y_train, Y_test = train_test_split(x, y, test_size=0.3,  
    ↪random_state = 42)  
print("size train : ",X_train.shape)  
print("size test : ",X_test.shape)
```

```
size train : (5534, 10)
```

```
size test : (2372, 10)
```

9 Setting up models

```
[21]: regression = [  
    Ridge(),  
    KNeighborsRegressor(),  
    LinearRegression(),  
    RandomForestRegressor(),  
    SVR(),  
    DecisionTreeRegressor(),  
    ElasticNet(),  
    Lasso(),  
    XGBRegressor()  
]
```

```
[22]: head = 10  
for model in regression[:head]:  
    start = time()  
    model.fit(X_train, Y_train)  
    train_time = time() - start  
    start = time()  
    Y_pred = model.predict(X_test)  
    predict_time = time()-start  
    print(model)  
    print("\t Temps d'entrainement : %0.3fs" % train_time)  
    print("\t Temps de prédiction : %0.3fs" % predict_time)  
    print("\t MAE score :", mean_absolute_error(Y_test, Y_pred))  
    print("\t R2 score :", r2_score(Y_test, Y_pred))  
    print("\t Max_error : ", max_error(Y_test, Y_pred))  
    print("\t MSE score : ", mean_squared_error(Y_test, Y_pred))  
    print()
```

```
Ridge()  
    Temps d'entrainement : 0.014s  
    Temps de prédiction : 0.001s  
    MAE score : 0.02996937997545131
```

R2 score : 0.563275406429672
Max_error : 0.4539752867762097
MSE score : 0.0029490586316302625

KNeighborsRegressor()

Temps d'entrainement : 0.015s
Temps de prédiction : 0.114s
MAE score : 0.009781175000509661
R2 score : 0.9246806113199936
Max_error : 0.27733196817131717
MSE score : 0.0005086072471897975

LinearRegression()

Temps d'entrainement : 0.018s
Temps de prédiction : 0.003s
MAE score : 0.03028970476108429
R2 score : 0.5578155727903901
Max_error : 0.5171904628972888
MSE score : 0.002985927106083595

RandomForestRegressor()

Temps d'entrainement : 1.698s
Temps de prédiction : 0.063s
MAE score : 0.007393877374907403
R2 score : 0.9686006258702937
Max_error : 0.1779911285187745
MSE score : 0.00021202972460969646

SVR()

Temps d'entrainement : 0.062s
Temps de prédiction : 0.015s
MAE score : 0.028460252361910517
R2 score : 0.7667076511856659
Max_error : 0.3561446033383627
MSE score : 0.0015753470839361351

DecisionTreeRegressor()

Temps d'entrainement : 0.021s
Temps de prédiction : 0.001s
MAE score : 0.008135735594842439
R2 score : 0.9587544824359456
Max_error : 0.2156469191928867
MSE score : 0.00027851751739908505

ElasticNet()

Temps d'entrainement : 0.009s
Temps de prédiction : 0.005s
MAE score : 0.043042230420574736

```
R2 score : -0.00010336126465948503
Max_error : 0.5888351724741655
MSE score : 0.006753371560663051
```

Lasso()

```
Temps d'entrainement : 0.008s
Temps de prédiction : 0.001s
MAE score : 0.043042230420574736
R2 score : -0.00010336126465948503
Max_error : 0.5888351724741655
MSE score : 0.006753371560663051
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)
Temps d'entrainement : 0.242s
Temps de prédiction : 0.005s
MAE score : 0.0069020062504309134
R2 score : 0.9733403469987905
Max_error : 0.17102317009828172
MSE score : 0.00018002393489393432
```

10 THE WINNER IS XGBRegressor AND RandomForestRegressor

10.1 looking for the best parameters

10.2 RandomForestRegressor

```
[23]: modelRFR= RandomForestRegressor()
      modelRFR
```

```
[23]: RandomForestRegressor()
```

```
[24]: modelRFR.fit(X_train, Y_train)
```

```
[24]: RandomForestRegressor()
```

```
[25]: parameters = {'n_estimators': np.arange(1,30), 'criterion': ["squared_error",  
↪ "friedman_mse", "absolute_error", "poisson"],  
                  'max_features': ["sqrt", "log2", "None"], 'random_state': np.  
↪ arange(1,5)}
```

```
[26]: kf = KFold(n_splits = 5, shuffle=True, random_state=5)  
grid = GridSearchCV(modelRFR, parameters, cv=kf, verbose=1)
```

```
[27]: grid.fit(X_train, Y_train)
```

Fitting 5 folds for each of 1392 candidates, totalling 6960 fits

```
[27]: GridSearchCV(cv=KFold(n_splits=5, random_state=5, shuffle=True),  
                  estimator=RandomForestRegressor(),  
                  param_grid={'criterion': ['squared_error', 'friedman_mse',  
                                           'absolute_error', 'poisson'],  
                              'max_features': ['sqrt', 'log2', 'None'],  
                              'n_estimators': array([ 1,  2,  3,  4,  5,  6,  7,  8,  
9, 10, 11, 12, 13, 14, 15, 16, 17,  
18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29])},  
                              'random_state': array([1, 2, 3, 4])},  
                  verbose=1)
```

```
[28]: grid.best_estimator_
```

```
[28]: RandomForestRegressor(criterion='friedman_mse', max_features='sqrt',  
                          n_estimators=26, random_state=4)
```

```
[29]: grid.best_score_
```

```
[29]: 0.9438128032607228
```

```
[30]: modelRFR = grid.best_estimator_  
modelRFR
```

```
[30]: RandomForestRegressor(criterion='friedman_mse', max_features='sqrt',  
                          n_estimators=26, random_state=4)
```

```
[31]: modelRFR.fit(X_train, Y_train)
```

```
[31]: RandomForestRegressor(criterion='friedman_mse', max_features='sqrt',  
                          n_estimators=26, random_state=4)
```

```
[32]: modelRFR.score(X_train, Y_train)
```

```
[32]: 0.9904613453056764
```


```
[33]: modelRFR.score(X_test, Y_test)
```

```
[33]: 0.9707851741448748
```

```
[34]: Y_pred_RFR = modelRFR.predict(X_test)
      print(Y_pred_RFR)
```

```
[0.05589856 0.0483477 0.01444343 ... 0.05063276 0.08351529 0.07728772]
```

```
[35]: plt.grid(True)
      sns.distplot(y, label='data real')
      sns.distplot(Y_pred_RFR, label='data predict')
      plt.legend()
      plt.show()
```



10.3 XGBRegression

```
[36]: modelXG = XGBRegressor()
      modelXG
```

```
[36]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytree=None, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric=None, feature_types=None,
```

```

gamma=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=None, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)

```

```
[37]: modelXG.fit(X_train, Y_train)
```

```

[37]: XGBRegressor(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=None, n_jobs=None,
    num_parallel_tree=None, random_state=None, ...)

```

```

[38]: parameters = {'n_estimators': np.arange(1,50), 'max_depth': np.arange(1,5),
    'max_features': ["sqrt", "log2", "None"], 'random_state': np.
    ↳arange(1,5)}

```

```

[39]: kf1 = KFold(n_splits = 5, shuffle=True, random_state=5)
    grid1 = GridSearchCV(modelXG, parameters, cv=kf1, verbose=1)

```

```
[40]: grid1.fit(X_train, Y_train)
```

Fitting 5 folds for each of 2352 candidates, totalling 11760 fits

```

[40]: GridSearchCV(cv=KFold(n_splits=5, random_state=5, shuffle=True),
    estimator=XGBRegressor(base_score=None, booster=None,
        callbacks=None, colsample_bylevel=None,
        colsample_bynode=None,
        colsample_bytree=None, device=None,
        early_stopping_rounds=None,
        enable_categorical=False, eval_metric=None,
        feature_types=None, gamma=None,
        grow_policy=None, importance_type=None,
        inter...
        multi_strategy=None, n_estimators=None,
        n_jobs=None, num_parallel_tree=None,
        random_state=None, ...),
    param_grid={'max_depth': array([1, 2, 3, 4]),
        'max_features': ['sqrt', 'log2', 'None'],

```

```

        'n_estimators': array([ 1,  2,  3,  4,  5,  6,  7,  8,
 9, 10, 11, 12, 13, 14, 15, 16, 17,
 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,
 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49]),
        'random_state': array([1, 2, 3, 4])},
        verbose=1)

```

```
[41]: modelXG = grid1.best_estimator_
      modelXG
```

```
[41]: XGBRegressor(base_score=None, booster=None, callbacks=None,
  colsample_bylevel=None, colsample_bynode=None,
  colsample_bytree=None, device=None, early_stopping_rounds=None,
  enable_categorical=False, eval_metric=None, feature_types=None,
  gamma=None, grow_policy=None, importance_type=None,
  interaction_constraints=None, learning_rate=None, max_bin=None,
  max_cat_threshold=None, max_cat_to_onehot=None,
  max_delta_step=None, max_depth=4, max_features='sqrt',
  max_leaves=None, min_child_weight=None, missing=nan,
  monotone_constraints=None, multi_strategy=None, n_estimators=49,
  n_jobs=None, num_parallel_tree=None, ...)

```

```
[42]: grid1.best_score_
```

```
[42]: 0.9325353333323065
```

```
[43]: grid1.best_estimator_
```

```
[43]: XGBRegressor(base_score=None, booster=None, callbacks=None,
  colsample_bylevel=None, colsample_bynode=None,
  colsample_bytree=None, device=None, early_stopping_rounds=None,
  enable_categorical=False, eval_metric=None, feature_types=None,
  gamma=None, grow_policy=None, importance_type=None,
  interaction_constraints=None, learning_rate=None, max_bin=None,
  max_cat_threshold=None, max_cat_to_onehot=None,
  max_delta_step=None, max_depth=4, max_features='sqrt',
  max_leaves=None, min_child_weight=None, missing=nan,
  monotone_constraints=None, multi_strategy=None, n_estimators=49,
  n_jobs=None, num_parallel_tree=None, ...)

```

```
[44]: modelXG.score(X_train, Y_train)
```

```
[44]: 0.9772679815268545
```


```
[45]: modelXG.score(X_test, Y_test)
```

```
[45]: 0.9613768098316412
```

```
[46]: Y_pred_XG = modelXG.predict(X_test)
print(Y_pred_XG)
```

```
[0.04059134 0.05420575 0.02151247 ... 0.05664548 0.08507872 0.10265834]
```

```
[47]: plt.grid(True)
sns.distplot(y, label='data real')
sns.distplot(Y_pred_XG, label='data predict')
plt.legend()
plt.show()
```



11 Conclusion

```
[48]: finale_report = pd.DataFrame({
    'Model': ['RandomForestRegression()', 'XGBRegression()'],
    'Score Train': [modelRFR.score(X_train, Y_train), modelXG.score(X_train, Y_train)],
    'Score Test': [modelRFR.score(X_test, Y_test), modelXG.score(X_test, Y_test)]
})
```

```
[49]: print(finale_report)
```


	Model	Score Train	Score Test
0	RandomForestRegression()	0.990461	0.970785
1	XGBRegression()	0.977268	0.961377

12 Conclusion

12.1 THE WINNER IS RANDOMFORESTREGRESSION